

Windows Internals: Advanced Exploit Mitigation

Fabian Nguyen

Abstract

We take a look at general techniques used by attackers to compromise Windows systems and some fundamental defense mechanisms against them. This paper will provide an overview of Microsoft's latest additions to the security concept of the Windows operating system (OS), analyze inherent flaws in their design and take a brief look at already existing attacks.

1 Introduction

In an increasingly digitalized world an overwhelming amount of private and/or safety-critical information and data is stored on computers. Windows is by far the most used operating system on desktop PCs [1] and therefore a main target of attackers to compromise data or computer systems. A common intent of attackers is to steal an individual's private/confidential data or install ransomware to extort money. Naturally, as the amount and complexity of attacks rise, OS vendors are forced to put an increasingly high amount of effort into mitigating existing weaknesses and deny attackers further possibilities to compromise the OS. Even though this is the case, the amount of potentially abusable vulnerabilities in Windows has been increasing, instead of decreasing.

Year	DoS	Code Execution	Overflow	Gain Information	Gain Privileges	Bypass Something
2015	4	19	6	5	26	10
2016	6	47	23	31	82	19
2017	32	50	16	108	19	18
2018	21	45	19	72	1	39
2019	28	124	101	73	2	10
Total	91	285	165	289	130	96
% of All	8.2	25.7	14.9	26.0	11.7	8.6

Figure 1: Amount of documented vulnerabilities in the Windows operating system (XSS with 0.3% share omitted for clarity) [2].

Note that the spike in *Gain Information* vulnerabilities in 2017 is inflated by a family of attacks widely known as Meltdown/Spectre.

In figure 1, we can see that the three largest groups of vulnerabilities consist of *Gain Information*, *Code Execution* and *Overflow*. These three categories are not entirely separated from each other. An attacker that is able to execute code on a machine often does so in order to gain information and overflows are often the reason why an attacker can execute code in the first place. We will take a look at how these issues play together and what has been done against them in the past before taking a deeper dive into Microsoft's newest defense additions.

2 Techniques and Data of Interest

More specifically, we will analyze Microsoft's two pillars of "Control-flow Enforcement" called Control Flow Guard (CFG) and Shadow Stack. Before doing so, we will take a look at why such sophisticated technologies are needed in the first place and what has been done in the interest of security before. Both of these technologies and relevant underlying concepts, namely overflows, ASLR, DEP, and ROP, will be explained, followed by a further going analysis of their strengths and weaknesses, and mentioning of existing attacks where applicable.

3 Overflows

An overflow happens when a program writes data to memory beyond the limits of the underlying data structure.




Figure 2: The password's length is too large to fit into the designated buffer. However, `strcpy()` will not check for valid sizes so it will continuously write beyond the buffer.

One common example of this is a stack buffer overflow caused by an incorrect use of the function `strcpy()`¹.

¹`char * strcpy (char * destination, const char * source)` is a C function that copies a string into the

More precisely, an overflow can occur when the given input is longer than the buffer one writes to, as seen in figure 2. Relatively small-in-size overflows are not always easy to spot and often remain unidentified if they do not cause immediate errors. In addition to causing faulty program behavior, this also provides a critical attack surface. To see why this is the case let us take a look at a typical stack layout with only one buffer present right at the beginning.

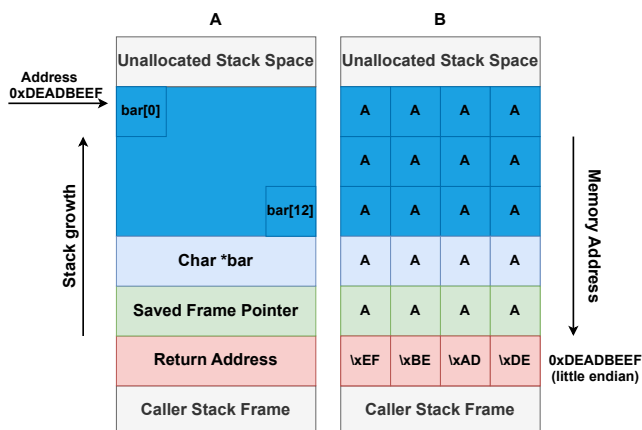


Figure 3: A: Typical Stack Layout after allocating a local buffer within a function with no arguments. B: The same stack after a misuse of `strcpy()` ^[3]

As we can see in Figure 3 (A) the stack will usually contain a return address right at the bottom, a frame pointer on top of it, and then local data that is used by the current function. One will also need to keep the parent function's data saved below (stack grows upwards in this example). There is only one character buffer of size 12 and a pointer to it. Suppose one wants to copy a user-input string into this buffer e.g., by using `strcpy()`. The user may now unknowingly or perhaps purposely overflow this buffer by entering a string that is longer than 12 characters. As we observed earlier this will result in a memory write beyond the buffer's bounds. Figure 3 (B) shows how the same stack might look after an overflow occurred. We can see that our pointer to the buffer is overwritten, as well as the frame pointer we had saved on the stack earlier. Most importantly though, the return address is also corrupted. In the given figure, the input is constructed so that the part that is written over the return address resembles an address itself. In this case, it is just the address of the buffer again. At some point, our given function will attempt to return to this address. However, once it does so, all it will be able to read from this

specified destination buffer

address is a string of 'A's which certainly does not resemble a valid sequence of code. The program will fault and terminate. This problem has existed for as long as the concept of the stack itself, so naturally techniques to (try) prevent this from happening were implemented long ago.

4 Data Execution Prevention

One intuitive way to address the unintended execution of (possibly injected) code on the stack is the use of Data Execution Prevention (DEP). DEP follows a very simple approach to prevent the execution of code from malicious addresses. Note that an attacker that uses stack overflows can usually only manipulate memory which is typically used for non-executable data. Based on this assumption, we could mark the entire memory area² (with an additional attribute for memory pages) that is used as a stack as *non-executable* or "NX" in short. Whenever an attempt to run code is made a check will ensure that the "NX"-Bit for the according page is not set. If it is, the program will terminate immediately. This approach offers an easy solution to prevent the *execution* of code on the stack. However, it does *not* prevent an attacker from redirecting control flow to already existing code. Attacks that abuse this flaw could therefore still compromise or even take over a program.³

5 Return-Oriented Programming

Return-oriented programming describes one of the most popular approaches to manipulating programs by hijacking control flow through the manipulation of return addresses. In essence, there are three steps to do:

1. Finding a security vulnerability that allows one to manipulate memory
2. Writing code (also called payload or shellcode)⁴ on the stack
3. Manipulating the return address to point to the injected code

Windows did not offer any kind of protection against ROP attacks until 2004 when DEP was introduced. However, as already mentioned, attackers quickly

²Sometimes code needs to be run directly from the stack e.g., just-in-time compiled JavaScript code, so this is a very simplified approach

³A well-known way to do this is a return-to-libc attack.

⁴The name shellcode comes from the fact that such attacks often inject code to open a shell on the target system

overcame this barrier by using already existent code which would not be marked non-executable [4]. A change to the calling convention that should have theoretically stopped this from happening was to require function arguments to be passed in registers instead [5]. Registers are not as easy to manipulate due to not being as easily writable per overflows, though attackers easily overcame this restriction as well. Instead of using complete functions they would now use only small portions of a function [6]. Naturally, instruction sequences that allowed to manipulate registers were especially useful, in order to invoke complete functions again. However, this is not needed as the usage of such so called "gadgets" is already turing-complete given a *big enough* [7] program. This is made even easier by the fact that one can use instruction sequences that were not supposed to be in the program to begin with. To see how this is possible, recall how machine code is written and read by the CPU. Note that in contrast to natural language, machine code does not include any white space (e.g., spaces or slashes) so one may start reading wherever they want. The following figure shows an example where two instructions are split into four just by removing one byte at the start.

```

1 F7 C7 07 00 00 00 - test $0x00000007, %edi
2 0F 95 45 C3       - setnzb -61(%ebp)
3
4 Missing the first Byte (F7) :
5
6 C7 07 00 00 00 0F - movl $0x0f000000, %edi
7 95                - xchg %ebp, %eax
8 45                - inc %ebp
9 C3                - ret

```

Figure 4: An example of an unintended use of machine code [8]. Ignoring the first byte results in a sequence of four instructions instead of two.

Note that the sequence we constructed ends with a ret instruction which allows us to continuously chain such gadgets. We can easily find such sequences by searching for ret instructions (in this example, the byte C3) and then going as far back as useful instructions follow [7].

6 Address Space Layout Randomization

Address Space Layout Randomization (ASLR) was implemented to address this issue. Since an attacker that uses preexisting code needs to be able to tell where the code they want to use is, an easy solution is to prevent them from addressing said code. ASLR does this by randomizing the arrangement of

segments like the stack, heap, and DLLs in memory. This means that using a fixed address to reference memory will most likely result in a segmentation error or the wrong code being targeted at best⁵. In theory, ASLR should prevent ROP attacks completely since no attacker should be able to locate any specific code to begin with. However, a single pointer leak may compromise the location of any of the segments and thus break ASLR completely. Consequently, vulnerabilities in both DEP and ASLR (even combined) have been found and successfully exploited [9]. How this was done will not be covered in this report, but it is important to see that these techniques were not sufficient protection after all.

7 Shadow Stack

The idea of a shadow stack follows a different approach to the issue. Instead of trying to prevent the manipulation of the return address, an additional check is built-in to make sure it was not altered. In order to do this, the shadow stack saves all return addresses in a separate memory location. Every call instruction in a program pushes the return address of the parent routine on the normal stack *and* on the shadow stack. Every return instruction pops one address off the shadow stack and compares it to the address on the normal stack.

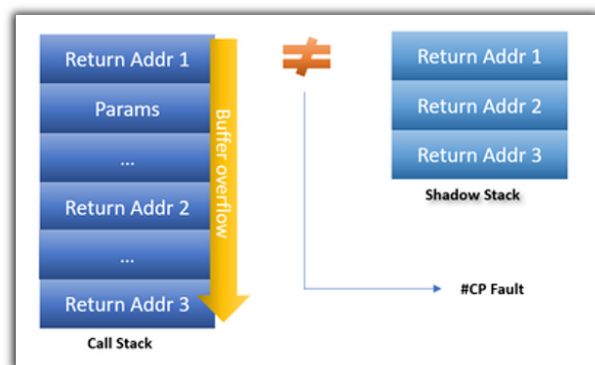


Figure 5: Shadow Stack Overview. Every return address on the stack is matched by a copy of that return address on the shadow stack.

In order to implement this check, a new register called shadow stack pointer(SSP) and an additional memory page attribute are used. Writes to the shadow stack are restricted to control transfer instructions (*call/ret*) and shadow stack man-

⁵32 bit Windows randomizes 8 of the address' bits, 64 Bit Windows can randomize a total of 19 bits.

agement instructions.⁶ Note that calls to exception handlers and *longjumps* may cause issues as calls and returns are not necessarily perfectly matched if they occur [10]. One solution to this is to continuously pop addresses off the shadow stack until a matching address is found. In this case, a program would be terminated if the shadow stack runs empty without match. Implementing special treatment for these troublesome calls instead could increase runtime performance but include a higher memory overhead and increase complexity. Using a parallel shadow stack removes the necessity of considering mismatches between call/ret instructions. Parallel shadow stacks are placed at a fixed offset relative to the main stack and all return addresses are placed parallel to the return addresses on the main stack (hence the name parallel shadow stack). Parallel shadow stacks have drastically reduced performance overhead (generally less than half) as it is no longer necessary to maintain the SSP (i.e. copying it to/from memory) but sacrifice some security to achieve this result [11]. Note that shadow stacks only provide integrity of "backward jumps" (also called backward edges). "Forward jumps" may still be manipulated while a shadow stack is being used.

7.1 Shadow Stack Revisited

One main problem with the shadow stack is that no one (except for researchers and employees that have early access to hardware that is still in development) is able to make use of it as of now. Since Microsoft's implementation of the Shadow Stack requires an additional register it is only usable with certain CPUs (more specifically those using a chipset that supports Intel Control-Flow-Enforcement technology) [12]. Intel's first generation of CPUs that use this chipset are yet to be released⁷ and so are AMD's.

ARM will not be releasing any compatible CPUs [13]. Additionally, once compatible CPUs are released it will still take a couple of years for them to be used widely. Considering that a software version of the shadow stack dubbed Return Flow Guard (RFG) [14] was already in development it is not quite apparent why this approach was not further pursued to bridge the time needed for hardware shadow stacks to spread. According to unofficial sources⁸ development of the RFG was discontin-

ued because a vulnerability (more specifically, a race condition) in the implementation was found. Performance overhead can range from as low as 3.5 (parallel) to roughly 10% (traditional shadow stacks) [11]. As Microsoft's implementation of the shadow stack will probably be neither parallel nor a basic traditional shadow stack we can expect the overhead to be anywhere in between these bounds. Note that the shadow stack only validates return addresses. Function parameters are (intentionally) not protected. Considering that it is unlikely for an attacker to be able to manipulate a function's parameters in between the call and the function running, this makes sense. Even if one were to try protecting them, it is unclear in which order the parameters should be saved to the shadow stack as the order of their use may be undetermined. It would also require doing this check for their first use only as they may be legitimately altered within the function. Lastly, according to Microsoft the use of shadow stacks will be opt-in "at first" per linker-flag for apps and DLLs for compatibility reasons. This will further increase the time needed for the shadow stack to spread.

8 Control Flow Guard

To take care of the Forward Edge the Control Flow Guard was shipped with Windows 8.1 Update 3. More specifically, CFG aims to prevent the misuse of indirect jumps. To do this developers can compile their programs with the CFG flag and ensure that their program will only run on memory areas that have been marked "safe" earlier. Here, "safe" means that the address of every indirect jump refers to a valid function in the program. In order to check this efficiently at run time, a bitmap of the starting addresses of every function is created at compile time where every bit in the bitmap corresponds to 8/16⁹ bytes in the address space. If a function starts within the block of addresses that correspond to a given bit it is set to '1', otherwise it is '0'. A guard instruction is then inserted before every indirect jump which validates that the matching bit for the target address in the bitmap is set to 1 (see figures 6 and 7). If it is not, the program is terminated immediately. If a CFG compatible program is run on a Windows version that does not support CFG the call simply does nothing. To increase performance, system processes which are run as "protected" processes use shared bitmaps for DLLs. This means that per DLL only one bitmap is needed for all protected processes unless a process tries to alter the bitmap. In this

⁶There is no point in having a separate stack to save data when its explicitly writable.

⁷Intel Tiger Lake CPUs are announced for late 2020

⁸<https://www.techrepublic.com/article/windows-10-security-how-the-shadow-stack-will-help-to-keep-the-hackers-at-bay>

⁹8 bytes per bit at first, later changed to tuple per 16 bytes

```

mov     ecx, 3E8h
rep stosd
mov     esi, [esi]
push    1
call    esi
add     esp, 4
xor     eax, eax

```

Pointer to fake object constructed by attacker

Call to the 1st stage shellcode

Figure 6: Example code without CFG ^[15]

```

mov     ecx, 3E8h
rep stosd
mov     esi, [esi]
mov     ecx, esi ; Target
push    1
call    @guard_check_icall@4 ; _guard_check_icall(x)
call    esi
add     esp, 4
xor     eax, eax

```

Figure 7: Example code with CFG enabled ^[15]

case, the bitmap is mapped into the process' private memory. Note that dynamically generated functions cannot be protected this way as explicitly allocated and *executable* marked memory has its corresponding bits in the bitmap implicitly set to 1. This is because dynamically allocated memory can not be considered at compile time.

8.1 Control Flow Guard revisited

One major flaw in the initial design of CFG lies in one of its main assumptions: That the addresses of functions are aligned to eight bytes. While this assumption holds true for the compiled program itself (since the compiler can enforce the alignment) it certainly is not for DLLs. To make up for this issue the CFG uses a mapping of two bits (a tuple) for every 16 bytes since Windows 10 instead where (0,0) means no function starts in the corresponding address block, (1,0) means a function starts at the start of the block and (1,1) means a function starts anywhere in the block. Note that this still provides a sort of margin for attacks when alignment is not enforced as an attacker can still manipulate a jump to point to anywhere within a (1,1) block (15 bytes of unintended viable targets in the worst case!). This is demonstrated in figure 8.

Another oversight that was found rather quickly: The guard function was not properly protected. In 2015 Zhang Yunhai showed that by using a (custom-)heap-allocation and freeing it afterwards one may make its memory area writable again under certain circumstances [16]. Yet another major issue that has not been fixed as of now: Lots of Windows system DLLs are not compiled to use CFG (145 in Windows 10 1511 / still 50+ in build 18362) [17]. Last but not least, note that any module, even if it is compiled to

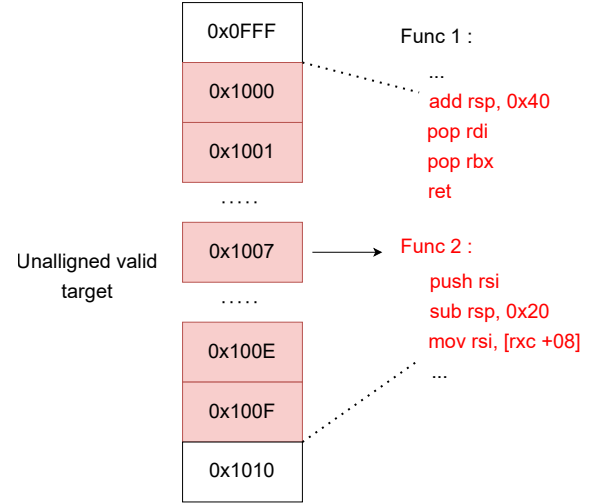


Figure 8: An example of why unaligned addresses can be a major vulnerability. Func2 makes the whole 16 byte range from 0x1000 to 0x100F valid, which contains instructions of func1's epilogue ^[15]

use CFG, is vulnerable when it is used by a program that is not, since CFG will be disabled in the entire process.

9 Approach

10 Conclusion

While both CFG and Shadow Stacks are valuable additions to the Windows security concept, Microsoft's war on hackers is far from over. As we have seen, CFG is vulnerable by design and has already been bypassed multiple times. The shadow stack has not been enabled for widespread use in Windows yet, but it is highly unlikely that no vulnerabilities will be found. Microsoft will have to continuously monitor and refine their implementations of these techniques to fix vulnerabilities and deny surface for attacks. Furthermore, performance (and possibly memory-) overhead may greatly impact the time needed for the shadow stack to be used widely.

References

- [1] Statcounter, "Desktop operating system market share worldwide."
- [2] <https://www.cvedetails.com/product/32238/Microsoft-Windows-10.html>, "Windows 10 security vulnerabilities."
- [3] Wikipedia, "Stack buffer overflow."

- [4] S. Designer, “Getting around non-executable stack.”
- [5] Microsoft, “x64 calling convention.”
- [6] S. Krahmer, “x86-64 buffer overflow exploits and the borrowed code chunks exploitation technique.”
- [7] A. Homescu, M. Stewart, P. Larsen, S. Brunthaler, and M. Franz, “Microgadgets: size does matter in turing-complete return-oriented programming,” in *Proceedings of the 6th USENIX conference on Offensive Technologies*, pp. 7–7, USENIX Association, 2012.
- [8] H. Shacham, “The geometry of innocent flesh on the bone: Return-into-libc without function calls (on the x86),” in *Proceedings of the 14th ACM conference on Computer and communications security*, pp. 552–561, 2007.
- [9] V. Katoch, “Bypassing aslr/dep.”
- [10] T. Zhang, “Shining light on shadow stacks.”
- [11] T. H. Dang, P. Maniatis, and D. Wagner, “The performance cost of shadow stacks and stack canaries,” in *Proceedings of the 10th ACM Symposium on Information, Computer and Communications Security*, pp. 555–566, 2015.
- [12] Intel, “Control-flow enforcement technology specification.”
- [13] M. Branscombe, “Windows 10 security: How the shadow stack will help to keep the hackers at bay.”
- [14] T. X. Lab, “Return flow guard.”
- [15] J. Tang and T. M. T. S. Team, “Exploring control flow guard in windows 10,” *Trend Micro Blog*, 2015.
- [16] Y. Zhang, “Bypass control flow guard comprehensively,” *Black Hat USA*, 2015.
- [17] [https://improsec.com/tech-blog/bypassing-control-flow-guard-on-windows-10-part ii](https://improsec.com/tech-blog/bypassing-control-flow-guard-on-windows-10-part-ii), “Bypassing cfg on windows 10.”